

An Analysis of Memory Management in HipHopVM

Benjamin Roberts, Nathan Yong, Jan Zimmer

HHVM Group, Summer Research Scholarship 2013/14

What is a HipHopVM?

Background

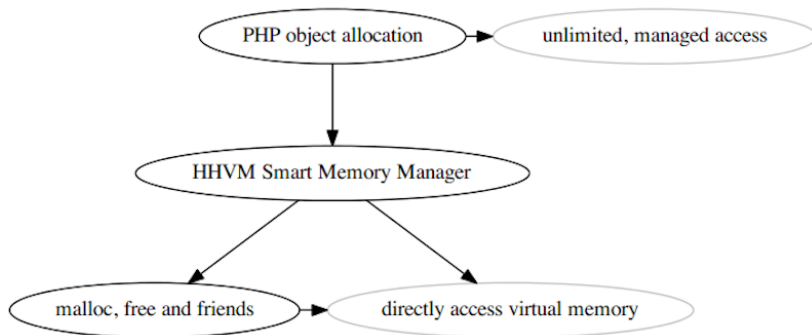
- HipHopVM is an Free and Open Source PHP engine
- Written primarily in C++ with moderate amounts of PHP and x86_64 assembly (1.2 million lines of code total)
- It uses a JIT compiler (though historically it translated PHP to C++ for AOT compilation)
- It is the PHP engine designed by and which hosts Facebook
- Its on GitHub! <https://github.com/facebook/hhvm>

Our Tasks

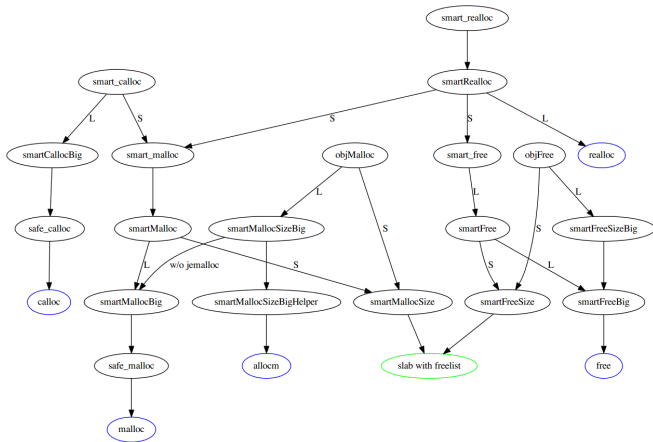
- To isolate the affect of naive reference counting on HHVM's performance
- To observe how memory access maps to actual physical memory access
- To map and analyse the behaviour of HHVM's internal memory management

Understanding the 'Smart' Memory Manager (Nathan Yong)

How do we view memory?



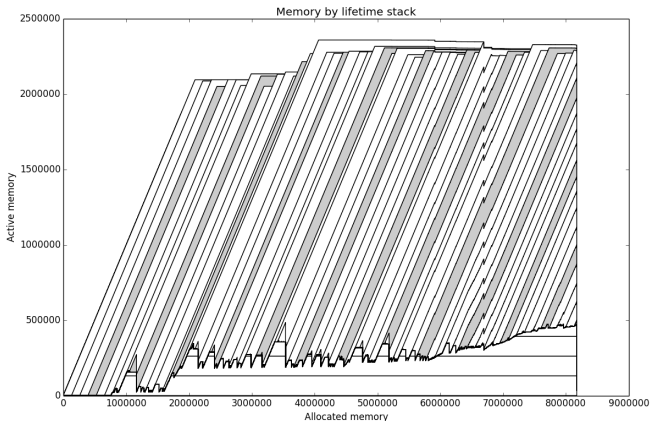
How HHVM sees memory




Why is this important?

- Smart memory manager not very well understood
- Usage of own tools very inconsistent


What is it doing?




What is it doing?



What is it doing?



What is it doing?



Where to from here?

- Opportunities for optimisation
- Unified consistent interface
- More efficient allocator design
- Possible use in conjunction with other tools

HHVM Without Reference Counting (Benjamin Roberts)

HHVM Without Reference Counting

- Who needs reference counts? (besides the PHP semantics)
- If we are freeing our heap at the end of request, why do we need immediate reclamation?
- Lets try nail down the performance penalty of immediate reference counting.

Reference Counting in the PHP language

- Based on explicit garbage collection (reference counting)
- Required by language semantics
- PHP is **pass by value**
- Pass by value can be slow due to large amount of copies (especially with large arrays)

- Solution: Copy on Write!
- Problem: Need to know current reference count
- Each mutation requires immediate increment and decrement of reference counts
- Advantage: Immediate garbage reclamation
- **This requirement will cause us major problems shortly**

Reference Counting in HHVM (C++)

- Reference counted objects have an `int_32t m_count` field and call a macro in `countable.h` containing various reference counting operations (`incRefCount()`, `hasMultipleRefs()`, `decRefAndRelease()` etc)
- Not consistently used (certain places directly mutate `m_count` or define separate ref-counting methods)
- Difficult to track down all mutations of `m_count`

Reference Counting in HHVM (JIT)

- `int_32t m_count` at a common offset (12 bytes) in all refcounted objects
- Several opcodes emit reference counting assembly via the JIT
- Intermediate Representation is optimised to remove inc/dec pairs where safe to do so
- Simpler to locate, more difficult to understand

Both these systems are implemented independent of each other (`m_count` and destructors shared) and the memory manager

Lets Modify HHVM

Removed/Disabled 3 forms of reference counting to create hhvmnocount:

- countable.h functions
- JIT emitted reference counting
- Various manual mutations and re-implementations (this took a while)

- Compared to regular HHVM, hhvmnocomp was much slower
- Completely different memory usage characteristics
- Haven't yet isolated the effects of reference counting

Continuous Allocation

In order to isolate the effects of reference counting a modified memory manager was used:

- Memory Manager no longer freed memory, No longer used free lists
- Large objects no longer special case
- Memory not freed at requests end
- Objects still '*swept*' due to unreclaimed File and Database objects causing errors.

- Resulted in the hhvmbump and hhvmbumpnocount branches.
- Reference counting was now isolated, time to run some tests.

Benchmark Setup

- Due to the nature of the changes, it was difficult to find representative PHP packages that ran on all builds
- A modified version of a benchmark included in HHVM (center-of-mass.php) was chosen due to its heavy memory usage
- center-of-mass.php is unlikely to represent the average PHP request but was the only test with sufficient memory usage for benchmarking

The configuration used for benchmarking:

- Linux kernel version: 3.12.6-300.fc20.x86_64
- CPU: Intel(R) Core(TM) i7-3770 CPU @ 3.40GHz
- Memory: 4x4G DDR3 memory at 1600MHz (no swap partition)
- internal ssd for HHVM builds
- Release configuration
- Apache Benchmark (ab) with various levels of concurrency and test lengths

Results:

- Results were graphed as surfaces using Matlab
- All results are graphed against total requests and concurrent-requests
- hhvmnocount is omitted from the results due to segmentation faults in Release configuration

Percentage Response Times (milliseconds, lower is better)




Figure : Time taken for 20% of requests to execute

- The lower 20% of response times are dependant on the build and the number of concurrent requests
- Small sample size, not very representative




Figure : Time taken for 50% of requests to execute

- Most interesting of the percentage response graphs
- Shows that hhvmbumpnocount performs the worst in majority of runs.
- This is contrary to expectations as it should have performed **less** operations than hhvmbump
- Will be discussed shortly




Figure : Time taken for 80% of requests to execute

- Large sample size now starting to include long-response times from warm-up period
- A lot of noise
- Still shows, like the previous graph, that hhvmbumpnocommit performs the worst.

Requests Processed per Second (higher is better)




Figure : Average requests per second of benchmark

- This graph shows that the removal of reference counting incurs a notable request processing penalty
- Penalty very noticeable in longer benchmark runs

Total Execution Time




Figure : Total execution time of benchmark

- Again shows that the removal of reference counting results in longer execution times

Why Did This Happen?

- Remains uncertain
- Benchmark chosen not representative of real PHP workload?
- Copy on Write behaviour?

Copy on Write

- As previously mentioned, copy on write requires exact reference counts
- `ArrayData` and `StringData` mutation behaviour based on the `hasMultipleRefs()` call (which is inaccurate in `hhvmbumpnocount` build)
- Over zealous copying may have occurred on mutation, resulting in performance penalty
- Could be confirmed by profiling and comparing memory usage of `hhvmbump` and `hhvmbumpnocount`

Further Work

Due to time constraints, several questions and problems remain unsolved:

- Identify source of negative result
- Re-run benchmark with Copy on Assignment semantics (potential method for previous point)
- Benchmark true request based GC (This was attempted early on before focus shifted to reference counting)
- Analyse the relationship between memory usage and response time (these modifications begin make memory a player in processing bottlenecks)
- Preserve copy-on-write behaviour without reference counting

Physical Memory Profile (Jan Zimmer)

Memory Utilisation

Memory

Memory utilisation



Memory utilisation



Memory utilisation




Memory utilisation



Memory utilisation




What HHVM can give us



What Valgrind can give us




What Valgrind and HHVM can give us




How they actually fit together




How they actually fit together



How they actually fit together



An example of the result



Physical Memory Profile

- Tool: Valgrind
- Get Memory Sectors from HHVM
- Monitor memory loads and stores from valgrind
- Valgrind alternates with HHVM in execution